

Chap 8. Hashing (1)

Contents

1. Introduction
2. Static Hashing
3. Dynamic Hashing

8.1 Introduction

ADT Dictionary is

objects: a collection of $n > 0$ pairs, each pair has a key and an associated item

functions:

for all $d \in \text{Dictionary}$, $item \in \text{Item}$, $k \in \text{Key}$, $n \in \text{integer}$

Dictionary Create(*max_size*) ::= create an empty dictionary.

Boolean IsEmpty(*d*, *n*) ::= **if** (*n* > 0) **return** FALSE
else return TRUE

Element Search(*d*, *k*) ::= **return** item with key *k*,
return NULL if no such element.

Element Delete(*d*, *k*) ::= delete and return item (if any) with key *k*;

void Insert(*d*, *item*, *k*) ::= insert *item* with key *k* into *d*.

ADT 5.3: Abstract data type *dictionary*

- Examples of dictionaries
 - Spelling checker
 - Thesaurus
 - Data dictionary
 - Symbol table

- In chapter 5, we learned how to build a BST for dictionaries
 - The worst-case time complexity for a BST is $O(n)$
 - The worst-case time complexity for a balanced BST can be $O(\log n)$
- In this chapter, We will learn a technique called *hashing* that enables us to perform *search*, *insert*, and *delete* in $O(1)$ expected time.
 - Static hashing
 - Dynamic hashing

8.2 Static Hashing

8.2.1 Hash Tables

- In *static hashing*, the dictionary pairs are stored in a table ht , called the ***hash table***.
- Hash table with ***b buckets*** and ***s slots***
 - buckets : $ht[0], \dots, ht[b-1]$
 - Each bucket is capable of holding s dictionary pairs.
- Hash function $h(k)$
 - determines the address of a pair whose key is k
 - maps keys into buckets
 - is an integer in the range 0 through $b-1$

- $h(k)$ is the hash or home address of k
 - Under the ideal condition, dictionary pairs are stored in their home buckets.
-
- **Definition:**
 - The *key density* of a hash table is the ratio n/T , where n is the number of pairs in the table and T is the total number of possible keys.
 - The *loading density* or *loading factor* of a hash table is $\alpha = n / (sb)$.

- Suppose
 - keys are at most six characters long
 - The first character must be a letter
 - The remaining characters can be letters or digits
- Then, the number of possible keys
 - $T = \sum_{i=0}^5 26 \times 36^i > 1.6 \times 10^9$
 - But most applications use only very small fraction of it
- Hash table also uses only small number of buckets
 - The hash function h maps several different keys into the same bucket

Terminologies

- Two keys k_1 and k_2 are said to be *synonyms* with respect to h if $h(k_1) = h(k_2)$
- *Collision* : occurs when the home bucket for a new pair is not empty at the time of insertion
- *Overflow* : hash a new identifier into a full bucket
- Collision and overflow occur at the same time if each bucket has 1 slot.

Example 8.1

- Suppose $b=26$, $s=2$, $n=10$
 - $\alpha = 10 / 52 = 0.19$
- identifiers: ‘acos’, ‘define’, ‘float’, ‘exp’, ‘char’, ‘atan’, ‘ceil’, ‘floor’, ‘clock’, ‘ctime’
- $h(x)$: the first character of x
- ‘acos’ and ‘atan’ are *synonym*
- ‘clock’: *overflow*

	Slot 0	Slot 1
0	acos	atan
1		
2	char	ceil
3	define	
4	exp	
5	float	floor
6		
...		
25		

Figure 8.1: Hash table with 26 buckets and two slots per bucket

- Insert, delete, find
 - $O(s)$ if no overflow occurs
 - $O(1)$ if no collision occurs
- But, collision occurs for most cases, since the ratio b/T is usually very small.
- Hash Table Issues
 - Choice of hash function
 - Overflow handling method
 - Size (number of buckets) of hash table

8.2.2 Hash Functions

- Requirements for hash function
 - Easy to compute
 - Minimize the number of collisions
 - Unbiased
 - uniform hashing function :
 - the probability of $h(k) = i$ is $1/b$ for all buckets i

8.2.2.1 Division

- The most widely used hash function in practice
- Assume that the keys are nonnegative
- The key is divided by some number D
 - $h(k) = k \% D$
 - bucket addresses : $0 \sim D-1$
 - At least $b = D$ buckets

- *The number of overflows* on real-world dictionaries *is critically dependent on the choice of D.*
 - The degree of bias decreases as the smallest prime factor of D increases.
- The relaxed requirement on D
 - use odd D & set $b = D$
 - *Array doubling* results in increasing the number of buckets (and hence the divisor D) from b to $2b + 1$.

8.2.2.2 Mid-Square

- Square the key
- Take an appropriate number of bits from the middle of square → use it as a key
- If r bits are used, then the table size = 2^r

ex)

$$\begin{array}{r} 10100 \\ 10100 \\ \hline 00\color{blue}{11001}0000 \end{array}$$

5 bits are used.

8.2.2.3 Folding

- The key k is partitioned into several parts with the same length and the partitions are added to obtain the hash address
- *Shift folding*: add all characters into one
 - Suppose $k = 12320324111220$
 - Partition: $(x_1: 123, x_2: 203, x_3: 241, x_4: 112, x_5: 20)$
 - $x_1 + x_2 + x_3 + x_4 + x_5 = 699$
- *Folding at the boundaries*: reverse every other partition before adding
 - reverse x_2, x_4 and add them
 - $x_1 + \text{reverse}(x_2) + x_3 + \text{reverse}(x_4) + x_5 = 897$

8.2.2.5 Converting Keys to Integers

- To use some of the hash functions, keys need to be converted to nonnegative integers.
- **Example 8.3:** *[Converting String to Integers]*
 - Two ways : Program8.1, Program 8.2

```
unsigned int stringToInt(char *key)
{ /* simple additive approach to create a natural number
   that is within the integer range */
    int number = 0;
    while (*key)
        number += *key++;
    return number;
}
```

Program 8.1: Converting a string into a non-negative integer

```
unsigned int stringToInt(char *key)
{ /* alternative additive approach to create a natural number
   that is within the integer range */
    int number = 0;
    while (*key)
    {
        number += *key++;
        if (*key) number += ((int) *key++) << 8;
    }
    return number;
}
```

Program 8.2: Alternative way to convert a string into a non-negative integer

This results in a larger range for the integer returned by the function.

8.2.3 Overflow Handling

8.2.3.1 Open Addressing

- Find the closest unfilled bucket when overflow occur.
- *Linear probing, quadratic probing, rehashing, random probing*

Example 8.4

- 13-bucket table with one slot per bucket
- hash value
 - By the scheme of Program 8.1 and division hash function

Identifier	Additive Transformation	x	Hash
for	$102 + 111 + 114$	327	2
do	$100 + 111$	211	3
while	$119 + 104 + 105 + 108 + 101$	537	4
if	$105 + 102$	207	12
else	$101 + 108 + 115 + 101$	425	9
function	$102 + 117 + 110 + 99 + 116 + 105 + 111 + 110$	870	12

[0]	function
[1]	
[2]	for
[3]	do
[4]	while
[5]	
[6]	
[7]	
[8]	
[9]	else
[10]	
[11]	
[12]	if


Using a circular rotation, the next available bucket is at $ht[0]$.

Linear probing

- (1) Compute $h(k)$.
- (2) Examine the hash table buckets in the order $ht[h(k)]$, $ht[(h(k) + 1) \% b]$, \dots , $ht[(h(k) + j) \% b]$ until one of the following happens:
 - (a) The bucket $ht[(h(k) + j) \% b]$ has a pair whose key is k ; in this case, the desired pair has been found.
 - (b) $ht[(h(k)+j)\%b]$ is empty; k is not in the table.
 - (c) We return to the starting position $ht[h(k)]$; the table is full and k is not in the table.

Linear probing

```
element* search(int k)
{ /* search the linear probing hash table ht (each bucket has
   exactly one slot) for k, if a pair with key k is found,
   return a pointer to this pair; otherwise, return NULL */
    int homeBucket, currentBucket;
    homeBucket = h(k); //0 ≤ h(k) < b
    for (currentBucket = homeBucket; ht[currentBucket]
          && ht[currentBucket]->key != k;) {
        currentBucket = (currentBucket + 1) % b;
                    /* treat the table as circular */
        if (currentBucket == homeBucket)
            return NULL; /* back to start point */
    }
    if (ht[currentBucket] && ht[currentBucket]->key == k)
        return ht[currentBucket];
    return NULL;
}
```



Program 8.3: Linear probing

Linear probing

- When linear probing is used to resolve overflows, keys tend to cluster together.
- Example :
 - Input sequence: **acos, atoi, char, define, exp, ceil, cos, float, atol, floor, ctime**
 - hash function uses the first character
 - What happens when we try to enter “atol” ?

The average number of comparisons : $41/11 = 3.73$

bucket	x	buckets searched
0	acos	1
1	atoi	2
2	char	1
3	define	1
4	exp	1
5	ceil	4
6	cos	5
7	float	3
8	atol	9
9	floor	5
10	ctime	9
...		
25		

Figure 8.4: Hash table with linear probing (26 buckets, one slot per bucket)

Linear probing

- With uniform hash function, the expected average number of key comparisons p is $(2-\alpha)/(2-2\alpha)$, where α is the loading density.
 - Figure 8.4 : $\alpha=11/26 = 0.42$, $p = 1.36$
- The worst-case number of comparisons: $O(n)$

Improvements

- Quadratic probing
 - Search $h(k)$, $(h(k) + i^2) \% b$, $(h(k) - i^2) \% b$ for $1 \leq i \leq (b - 1)/2$, where b is the # of buckets
 - When b is a prime number of the form $4j+3$, where j is an integer, every buckets are examined

Prime	j	Prime	j
3	0	43	10
7	1	59	14
11	2	127	31
19	4	251	62
23	5	503	125
31	7	1019	254

Figure 8.5: Some primes of the form $4j + 3$

Improvements

- Rehashing
 - applying a series of hash functions h_1, h_2, \dots, h_m to reduce clustering
 - examine buckets using $h_i(k)$ ($1 \leq i \leq m$) sequentially

Improvements

- Random Probing ※ Exercises 5
 - The search for a key, k , in a hash table with b buckets is carried out by examining the buckets in the order $h(k)$, $(h(k)+s(i))\%b$, $1 \leq i \leq b-1$ where $s(i)$ is a pseudo random number.
 - *The random number generator* must satisfy the property that *every number from 1 to b-1* must be generated *exactly once* as i ranges from 1 to $b-1$

Example

Input sequence : 5 8 13 7 21 23

Random numbers : 5 2 3 7 1 4 6

Hash table : 8 buckets with 1 slot

$$k=5 : h(k) = 5 \% 8 = 5$$

$$k=8 : h(k) = 8 \% 8 = 0$$

$$k=13 : h(k) = 13 \% 8 = 5$$

$$(h(k)+s(1)) \% 8 = (5+5) \% 8 = 2$$

$$k=7 : h(k) = 7 \% 8 = 7$$

$$k=21 : h(k) = 21 \% 8 = 5$$

$$(h(k)+s(1)) \% 8 = (5+5) \% 8 = 2$$

$$(h(k)+s(2)) \% 8 = (5+2) \% 8 = 7$$

$$(h(k)+s(3)) \% 8 = (5+3) \% 8 = 0$$

$$(h(k)+s(4)) \% 8 = (5+7) \% 8 = 4$$

$$k=23 : h(k) = 23 \% 8 = 7$$

$$(h(k)+s(1)) \% 8 = (7+5) \% 8 = 4$$

$$(h(k)+s(2)) \% 8 = (7+2) \% 8 = 1$$

[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]
8	23	13		21	5		7

8.2.3.2 Chaining

- Maintain a linked list of synonyms for each bucket
 - s is flexible
 - Each list contains all the synonyms for that bucket

[0] → **acos atoi atol**
[1] → *NULL*
[2] → **char ceil cos ctime**
[3] → **define**
[4] → **exp**
[5] → **float floor**
[6] → *NULL*
...
[25] → *NULL*

The average number of
comparisons : $21/11 = 1.91$


Figure 8.6: Hash chains corresponding to Figure 8.4

```

element* search(int k)
/* search the chained hash table ht for k, if a pair with
this key is found, return a pointer to this pair;
otherwise, return NULL.
nodePointer current;
int homeBucket = h(k); //0 ≤ h(k) < b
/* search the chain ht[homeBucket] */
for (current = ht[homeBucket]; current;
                                current = current→link)
    if (current→data.key == k) return &current→data;
return NULL;
}

```

Program 8.4: Chain search



- With uniform hash function, the expected average number of key comparisons p is $\approx 1 + \alpha/2$.
 - Figure 8.6 : $\alpha = 11/26 = 0.42$, $p = 1.21$
- The worst-case number of comparisons: $O(n)$
 - can be reduced to $O(\log n)$ by sorting synonyms in a balanced search tree